


Experiment 5: Downstream Agent Evaluation

Held-out API (Zoom): base vs.
fine-tuned on EnterpriseSynth data



LoRA fine-tuning loss
(Qwen2.5-0.5B, 45 examples)

